



Relational Algebra

MSCI346
Lukasz Golab

Acknowledgement: slides derived from material provided by
Sliberschatz, Korth and Sudarshan, copyright 2019, www.db-book.com,
and by Prof. Wojciech Golab, ECE356



Learning Outcomes

- Develop a working knowledge of relational algebra:
 - basic operators
 - ▶ select, project, union, difference, Cartesian product, rename
 - additional operators
 - ▶ intersection, assignment, natural join
 - formula syntax
- Understand null values and 3-valued logic

- Textbook sections (6th ed.): 6.1, 6.2, 6.3



Relational Algebra

- Procedural language (as opposed to declarative).
- Six basic operators:
 - select: σ
 - project: Π
 - union: \cup
 - set difference: $-$
 - Cartesian product: \times
 - rename: ρ
- These operators take one or two relations as inputs and output a single relation.



Select Operation – Example

- Relation r :

A	B	C	D
α	α	1	7
α	β	5	7
β	β	12	3
β	β	23	10

- $\sigma_{A=B \wedge D > 5}(r)$

A	B	C	D
α	α	1	7
β	β	23	10



Select Operation

- Notation: $\sigma_p(r)$
- p is called the **selection predicate**
- Defined as:

$$\sigma_p(r) = \{t \mid t \in r \text{ and } p(t)\}$$

Where p is a formula in propositional logic consisting of **terms** connected by : \wedge (**and**), \vee (**or**), \neg (**not**)
Each **term** is one of:

$\langle \text{attribute} \rangle \quad op \quad \langle \text{attribute} \rangle \text{ or } \langle \text{constant} \rangle$

where op is one of: $=, \neq, >, \geq, <, \leq$

- Example of selection:

$$\sigma_{dept_name="Physics"}(instructor)$$



Project Operation – Example

- Relation r :

A	B	C
α	10	1
α	20	1
β	30	1
β	40	2

- $\Pi_{A,C}(r)$

A	C
α	1
α	1
β	1
β	2

 $=$

A	C
α	1
β	1
β	2



Project Operation

- Notation:

$$\Pi_{A_1, A_2, \dots, A_k}(r)$$

where A_1, A_2 are attribute names and r is a relation name.

- The result is defined as the relation of k columns obtained by erasing the columns that are not listed.
- Note: since relations are sets, duplicate rows “removed” from result.
- Example: To eliminate the *dept_name* attribute of *instructor*

$$\Pi_{ID, name, salary}(instructor)$$



Union Operation – Example

- Relations r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

- $r \cup s$:

A	B
α	1
α	2
β	1
β	3



Union Operation

- Notation: $r \cup s$
- Defined as:

$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$

- For $r \cup s$ to be valid.
 1. r, s must have the *same* **arity** (same number of attributes)
 2. The attribute domains must be **compatible**
(Example: 2nd column of r deals with the same type of values as does the 2nd column of s . Column names may be different.)
- Example: to find all courses taught in the Fall 2009 semester, or in the Spring 2010 semester, or in both

$$\Pi_{course_id}(\sigma_{semester="Fall" \wedge year=2009}(section)) \cup \Pi_{course_id}(\sigma_{semester="Spring" \wedge year=2010}(section))$$



Set difference of two relations

- Relations r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

- $r - s$:

A	B
α	1
β	1



Set Difference Operation

- Notation $r - s$
- Defined as:

$$r - s = \{t \mid t \in r \text{ and } t \notin s\}$$

- Set differences must be taken between **compatible** relations, just like unions.
- Example: to find all courses taught in the Fall 2009 semester, but not in the Spring 2010 semester

$$\Pi_{course_id}(\sigma_{semester="Fall" \wedge year=2009}(section)) - \Pi_{course_id}(\sigma_{semester="Spring" \wedge year=2010}(section))$$



Cartesian-Product Operation – Example

- Relations r, s :

A	B
α	1
β	2

r

C	D	E
α	10	a
β	10	a
β	20	b
γ	10	b

s

- $r \times s$:

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b



Cartesian-Product Operation

- Notation: $r \times s$
- Defined as:

$$r \times s = \{t q \mid t \in r \textbf{ and } q \in s\}$$

Note: “ $t q$ ” denotes a tuple obtained by concatenating together t and q .

- Note: the definition assumes that attributes of $r(R)$ and $s(S)$ are disjoint. (That is, $R \cap S = \emptyset$.)
- If attributes of $r(R)$ and $s(S)$ are not disjoint, then renaming must be used.



Composition of Operations

- Can build expressions using multiple operations
- Example: $\sigma_{A=C}(r \times s)$

- $r \times s$

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

- $\sigma_{A=C}(r \times s)$

A	B	C	D	E
α	1	α	10	a
β	2	β	10	a
β	2	β	20	b



Rename Operation

- Allows us to name, and therefore to refer to, the results of relational-algebra expressions.
- Allows us to refer to a relation by more than one name.
- Example:

$$\rho_x(E)$$

returns the expression E under the name X

- If a relational-algebra expression E has arity n , then

$$\rho_{x(A_1, A_2, \dots, A_n)}(E)$$

returns the result of expression E under the name X , and with the attributes renamed to A_1, A_2, \dots, A_n .



Example Query

- Find the largest salary in the university
 - Step 1: find instructor salaries that are less than some other instructor salary (i.e. not maximum)
 - using a copy of *instructor* under a new name *d*
 - ▶ $\Pi_{instructor.salary} (\sigma_{instructor.salary < d.salary} (instructor \times \rho_d (instructor)))$
 - Step 2: find the largest salary
 - ▶ $\Pi_{salary} (instructor) - \Pi_{instructor.salary} (\sigma_{instructor.salary < d.salary} (instructor \times \rho_d (instructor)))$



Formal Definition of Relational Algebra

- A **basic expression** in the relational algebra consists of either a relation in the database (e.g., *instructor*) or a constant relation (e.g., $\{(1, \text{Einstein}), (2, \text{Crick})\}$).
- A general **relational algebra expression** is either a basic expression or an expression constructed recursively using one of the following rules, where E_1 and E_2 denote existing relational-algebra expressions:
 - $E_1 \cup E_2$
 - $E_1 - E_2$
 - $E_1 \times E_2$
 - $\sigma_P(E_1)$, where P is a predicate on attributes in E_1
 - $\Pi_S(E_1)$, where S is a list comprising a subset of the attributes in E_1
 - $\rho_{x(A_1, A_2, \dots, A_n)}(E_1)$, where $x(A_1, A_2, \dots, A_n)$ is the new name for E_1 and its attributes



Additional Operations

For convenience, additional relational operators can be defined that do not add any expressive power to the relational algebra but simplify common queries.

- set intersection: \cap
- natural join: \bowtie
- theta join: \bowtie_{θ}
- assignment: \leftarrow
- set division: \div
- outer join (see textbook, covered in a later lecture on SQL)



Set-Intersection Operation

- Notation: $r \cap s$
- Defined as:

$$r \cap s = \{ t \mid t \in r \textbf{ and } t \in s \}$$

- Assume:
 - r, s have the *same arity*
 - attributes of r and s are compatible
- Note: $r \cap s = r - (r - s)$



Set-Intersection Operation – Example

- Relation r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

- $r \cap s$

A	B
α	2



Assignment Operation

- The assignment operation (\leftarrow) provides a convenient way to express complex queries.
 - Write query as a sequential program consisting of
 - ▶ a series of assignments
 - ▶ followed by an expression whose value is displayed as a result of the query.
 - Assignment must always be made to a temporary relation variable.
- Example: find the largest salary in the university (in two lines of “code”)

$$temp \leftarrow \Pi_{instructor.salary} (\sigma_{instructor.salary < d.salary} (instructor \times \rho_d (instructor)))$$
$$\Pi_{salary} (instructor) - temp$$



Natural-Join Operation

- Notation: $r \bowtie s$
- Let r and s be relations on schemas R and S respectively.
Then, $r \bowtie s$ is a relation on schema $R \cup S$ obtained as follows:
 - Consider each pair of tuples t_r from r and t_s from s .
 - If t_r and t_s have the same value on each of the attributes in $R \cap S$, add a tuple t to the result, where
 - ▶ t has the same value as t_r for attributes in R
 - ▶ t has the same value as t_s for attributes in S

- Example:

$R = (A, B, C, D)$

$S = (E, B, D)$

- Result schema = (A, B, C, D, E)
- $r \bowtie s$ is defined as:

$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B = s.B \wedge r.D = s.D} (r \times s))$$



Natural Join Example

- Relations r , s :

A	B	C	D
α	1	α	a
β	2	γ	a
γ	4	β	b
α	1	γ	a
δ	2	β	b

r

B	D	E
1	a	α
3	a	β
1	a	γ
2	b	δ
3	b	ϵ

s

- $r \bowtie s$

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ



Natural Join and Theta Join

- Find the names of all instructors in the Comp. Sci. department together with the course titles of all the courses that the instructors teach
 - $\Pi_{name, title} (\sigma_{dept_name="Comp. Sci."} (instructor \bowtie teaches \bowtie course))$
- Natural join is associative
 - $(instructor \bowtie teaches) \bowtie course$ is equivalent to $instructor \bowtie (teaches \bowtie course)$
- The **theta join** operation $r \bowtie_{\theta} s$ is defined as
 - $r \bowtie_{\theta} s = \sigma_{\theta} (r \times s)$
 - Example: $instructor \bowtie_{instructor.ID = teaches.ID} teaches$

(How is this different from $instructor \bowtie teaches$?)



Natural Join and Theta Join (Cont.)

- Example schema:

pet (*p_id*, *name*, *species*)

owner (*o_id*, *name*, *address*)

owns (*p_id*, *o_id*)

- Query: find the name and address of every dog that has an owner.
- Answer:

$temp \leftarrow (pet \bowtie owns)$

$temp2 \leftarrow temp \bowtie_{temp.o_id = owner.o_id} owner$

$\Pi_{temp.name, address} (\sigma_{species = \text{"dog"}} (temp2))$



Division Operator

- Given relations $r(R)$ and $s(S)$, such that $S \subset R$, $r \div s$ is the largest relation $t(R - S)$ such that

$$t \times s \subseteq r$$

- Example:

- let $r(ID, course_id) = \Pi_{ID, course_id}(takes)$ and
 $s(course_id) = \Pi_{course_id}(\sigma_{dept_name="Biology"}(course))$
- then $r \div s$ gives us students who have taken all courses in the Biology department

- Can rewrite $r \div s$ as

$$temp1 \leftarrow \Pi_{R-S}(r)$$

$$temp2 \leftarrow \Pi_{R-S}((temp1 \times s) - r)$$

$$temp1 - temp2$$



Division Operator (Cont.)

■ Example with real data:

- student IDs are 1, 2, 3, 4
- course IDs are 'BIO-101', 'BIO-301', and 'CS-101'
- $r = \{ (1, 'BIO-101'), (1, 'BIO-301'), (2, 'CS-301') \}$
- $s = \{ ('BIO-101'), ('BIO-301') \}$
- $r \div s = \{ (1) \}$

■ The three-line query for $r \div s$ can be executed as follows:

$$temp1 = \Pi_{R-S}(r) = \{ (1), (2) \}$$

$$temp1 \times s = \{ (1, 'BIO-101'), (1, 'BIO-301'), (2, 'BIO-101'), (2, 'BIO-301') \}$$

$$(temp1 \times s) - r = \{ (2, 'BIO-101'), (2, 'BIO-301') \}$$

$$\Pi_{R-S}((temp1 \times s) - r) = \{ (2) \}$$

$$temp1 - temp2 = \{ (1) \}$$



Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*.
- Aggregate functions (not covered in this lecture) simply ignore null values (as in SQL).
- For duplicate elimination and grouping (not covered in this lecture), *null* is treated like any other value, and two *null* values are assumed to be the same (as in SQL).